

Lecture3a

September 14, 2025

1 CS5489 - Machine Learning

2 Lecture 3a - Discriminative Classifiers

2.0.1 Dept. of Computer Science, City University of Hong Kong

3 Outline

1. Discriminative linear classifiers
2. Logistic regression
3. SVM & Kernel SVM
4. Summary

```
[1]: # setup
%matplotlib inline
import matplotlib_inline # setup output image format
matplotlib_inline.backend_inline.set_matplotlib_formats('retina')
import matplotlib.pyplot as plt
plt.rcParams['figure.dpi'] = 100 # display larger images
import matplotlib
from numpy import *
from sklearn import *
from scipy import stats
```

4 Classification with Generative Model

- Steps to build a classifier
 1. Collect training data (features $\mathbf{x}$ and class labels y)
 2. Learn class-conditional distribution (CCD), $p(\mathbf{x}|y)$.
 3. Use Bayes' rule to calculate class probability, $p(y|\mathbf{x})$.
- **Note:** the data is used to learn the CCD – the classifier is secondary.
 - Density estimation is an “ill-posed” problem – which density to use? how much data is needed?
- Advice from Vladimir Vapnik (inventor of SVM): > When solving a problem, try to avoid solving a more general problem as an intermediate step.
- **Discriminative solution**
 - Solve for the classifier $p(y|\mathbf{x})$ directly!

- Terminology
 - “**Discriminative**” - learn to directly discriminate the classes apart using the features.
 - “**Generative**” - learn model of how the features are generated from different classes.

5 Revisit the Naive Bayes Gaussian Classifier

- CCDs: assume the same variance for all Gaussians:
 - $p(\mathbf{x}|y=1) = \prod_{i=1}^D N(x_i|\mu_i, \sigma^2)$
 - $p(\mathbf{x}|y=2) = \prod_{i=1}^D N(x_i|\nu_i, \sigma^2)$
- prior:
 - $p(y=1) = \pi_1, p(y=2) = \pi_2$.
- look at the log-ratio of CCDs,

$$\begin{aligned}
 \log \frac{p(\mathbf{x}|y=1)}{p(\mathbf{x}|y=2)} &= \log \frac{\prod_{i=1}^D N(x_i|\mu_i, \sigma^2)}{\prod_{i=1}^D N(x_i|\nu_i, \sigma^2)} \\
 &= \sum_{i=1}^D \log N(x_i|\mu_i, \sigma^2) - \log N(x_i|\nu_i, \sigma^2) \\
 &= \sum_{i=1}^D -\frac{1}{2\sigma^2}(x_i - \mu_i)^2 + \frac{1}{2\sigma^2}(x_i - \nu_i)^2 \\
 &= \frac{1}{2\sigma^2} \sum_{i=1}^D (2x_i\mu_i - \mu_i^2 - 2x_i\nu_i + \nu_i^2) \\
 &= \frac{1}{2\sigma^2} \sum_{i=1}^D 2(\mu_i - \nu_i)x_i - \mu_i^2 + \nu_i^2
 \end{aligned}$$

- Thus

$$\log \frac{p(\mathbf{x}|y=1)}{p(\mathbf{x}|y=2)} = \frac{1}{\sigma^2} \sum_{i=1}^D (\mu_i - \nu_i)x_i + \frac{1}{2\sigma^2} \sum_{i=1}^D (\nu_i^2 - \mu_i^2)$$

- **Bayes decision rule:** Compute the posterior probability of each class $p(y=j|\mathbf{x})$
 - select class 1 when:

$$\begin{aligned}
 \log p(y=1|\mathbf{x}) &> \log p(y=2|\mathbf{x}) \\
 \log p(\mathbf{x}|y=1) + \log p(y=1) &> \log p(\mathbf{x}|y=2) + \log p(y=2) \\
 \log \frac{p(\mathbf{x}|y=1)}{p(\mathbf{x}|y=2)} + \log \frac{p(y=1)}{p(y=2)} &> 0
 \end{aligned}$$

- substituting for the CCDs and priors, the BDR is:
 - select class $y=1$ when:

$$\sum_{i=1}^D \frac{1}{\sigma^2} (\mu_i - \nu_i)x_i + \frac{1}{2\sigma^2} \sum_{i=1}^D (\nu_i^2 - \mu_i^2) + \log \frac{\pi_1}{\pi_2} > 0$$

```
[2]: # load iris data each row is (petal length, sepal width, class)
irisdata = loadtxt('iris2.csv', delimiter=',', skiprows=1)
```

```
X = irisdata[:,0:2] # the first two columns are features (petal length, sepal
    ↪width)
Y = irisdata[:,2]   # the third column is the class label (versicolor=1,
    ↪virginica=2)

print(X.shape)
```

(100, 2)

```
[3]: # get the NB Gaussian model from sklearn
model = naive_bayes.GaussianNB()

# fit the model
model.fit(X, Y)
v = mean(model.var_.ravel()) # make it shared variance
model.var_[:] = v
```

```
[4]: # a colormap for making the scatter plot: class 1 will be red, class 2 will be
    ↪green
mycmap = matplotlib.colors.LinearSegmentedColormap.from_list('mycmap',
    ↪["#FF0000", "#FFFFFF", "#00FF00"])

axbox = [2.5, 7, 1.5, 4] # common axis range

# a function for setting a common plot
def irisaxis():
    plt.xlabel('petal length'); plt.ylabel('sepal width')
    plt.axis([2.5, 7, 1.5, 4]); plt.grid(True)
```

```
[5]: ## Visualization
    ↪#####
def plot_ellipse(ax, musigma, color="k", lw=1):
    """
    Based on
    http://stackoverflow.com/questions/17952171/
    ↪not-sure-how-to-fit-data-with-a-gaussian-python.
    """

    mu, sigma = musigma

    # Compute eigenvalues and associated eigenvectors
    vals, vecs = linalg.eigh(sigma)

    # Compute "tilt" of ellipse using first eigenvector
    x, y = vecs[:, 0]
    theta = degrees(arctan2(y, x))
```

```

# Eigenvalues give length of ellipse along each eigenvector
# plot 2 stdevs
w, h = 2 * sqrt(vals) * 2

#ax.tick_params(axis='both', which='major', labelsize=20)
ellipse = matplotlib.patches.Ellipse(mu, w, h, angle=theta, fill=False,
↪color=color, lw=lw) # color="k")
ellipse.set_clip_box(ax.bbox)
ellipse.set_alpha(1.0)
ax.add_artist(ellipse)

```

```

[6]: def plot_posterior(model, axbox, mycmap, showlabels=True):
    xr = [ linspace(axbox[0], axbox[1], 200),
           linspace(axbox[2], axbox[3], 200) ]

    # make a grid for calculating the posterior,
    # then form into a big [N,2] matrix
    xgrid0, xgrid1 = meshgrid(xr[0], xr[1])
    allpts = c_[xgrid0.ravel(), xgrid1.ravel()]

    # calculate the posterior probability
    post = model.predict_proba(allpts)
    # extract the posterior for class 2, and reshape into a grid
    post1 = post[:,1].reshape(xgrid0.shape)

    # contour plot of the posterior and decision boundary
    plt.imshow(post1, origin='lower', extent=axbox, alpha=0.50, cmap=mycmap)
    if showlabels:
        plt.colorbar(shrink=0.6)
        CS = plt.contour(xr[0], xr[1], post1, cmap=mycmap, levels=[0.1, 0.3, 0.7, 0.
↪9], alpha=0.8)
        if showlabels:
            #plt.colorbar(CS)
            plt.clabel(CS, inline=1, fontsize=10)
        plt.contour(xr[0], xr[1], post1, levels=[0.5], linestyle='dashed',
↪colors='black')
    irisaxis()

```

```


[7]: pfig = plt.figure(figsize=(9,8))
plot_posterior(model, axbox, mycmap, showlabels=True)
plot_ellipse(plt.gca(), (model.theta_[0,:], diag(model.var_[0,:])), color='r')
plot_ellipse(plt.gca(), (model.theta_[1,:], diag(model.var_[1,:])), color='g')
plt.scatter(X[:,0], X[:,1], c=Y, cmap=mycmap, edgecolors='k')
plt.title('posterior p(y=2|x) with training data');
plt.close()

```

- Example

```
[8]: pfig
```

```
[8]:
```



- BDR in this case is a **linear** function
 - select class $y = 1$ when:

$$\sum_{i=1}^D \underbrace{\frac{1}{\sigma^2}(\mu_i - \nu_i)}_{w_i} x_i + \underbrace{\frac{1}{2\sigma^2} \sum_{i=1}^D (\nu_i^2 - \mu_i^2) + \log \frac{\pi_1}{\pi_2}}_b > 0$$

- * w_i is a per-feature weight
- * b is a bias term

- the BDR in this case is a **linear** classifier:
 - select class $y = 1$ when
 - * $\sum_{i=1}^D w_i x_i + b > 0$
 - * equivalently, $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b > 0$
- Here we obtain the weights $\mathbf{w}$ by learning the CCDs
 - assuming Naive Bayes Gaussians with shared variance.
 - this is a generative model since we learn how the data is generated for each class (CCDs).
- How to learn the linear classifier in a discriminative way?
 - directly learn the posterior $p(y|\mathbf{x})$.
 - we will look at a generic linear classifier.

6 Linear Classifier

- **Setup**
 - Observation (feature vectors) $\mathbf{x} \in \mathbb{R}^d$

- Class $y \in \{-1, +1\}$
- **Goal:** given a feature vector $\mathbf{x}$, predict its class y .
 - Calculate a *linear function* of the feature vector $\mathbf{x}$.
 - * $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b = \sum_{j=1}^d w_j x_j + b$
 - $\mathbf{w} \in \mathbb{R}^d$ are the weights of the linear function.
 - multiply each feature value with a weight, and then add together.
 - Predict from the value:
 - * if $f(\mathbf{x}) > 0$ then predict Class $y = 1$
 - * if $f(\mathbf{x}) < 0$ then predict Class $y = -1$
 - * Equivalently, $y = \text{sign}(f(\mathbf{x}))$

7 Geometric Interpretation

- The linear classifier separates the features space into 2 *half-spaces*
 - corresponding to feature values belonging to Class +1 and Class -1
 - the class boundary is normal to $\mathbf{w}$.
 - * also called the *separating hyperplane*.
- Example:

$$\mathbf{w} = \begin{bmatrix} 2 \\ 1 \end{bmatrix}, b = 0$$

```
[9]: def drawplane(w, b, wlabel=None, poscol=None, negcol=None):
    # w^T x + b = 0
    # w0 x0 + w1 x1 + b = 0
    # x1 = -w0/w1 x0 - b / w1

    # the line
    x0 = linspace(-10,10)
    x1 = -w[0]/w[1] * x0 - b / w[1]

    # fill positive half-space or neg space
    if (poscol):
        polyx = [x0[0], x0[-1], x0[-1], x0[0]]
        polyy = [x1[0], x1[-1], x1[0], x1[0]]
        plt.fill(polyx, polyy, poscol, alpha=0.2)

    if (negcol):
        polyx = [x0[0], x0[-1], x0[0], x0[0]]
        polyy = [x1[0], x1[-1], x1[-1], x1[0]]
        plt.fill(polyx, polyy, negcol, alpha=0.2)

    # plot line
    plt.plot(x0, x1, 'k-', lw=2)

    # the w
```

```

    if (wlabel):
        xp = array([0, -b/w[1]])
        xpw = xp+w
        plt.arrow(xp[0], xp[1], w[0], w[1], width=0.01, head_width=0.3,
        ↪fc='black')
        plt.text(xpw[0]-0.5, xpw[1], wlabel)

```

```

[10]: linclass = plt.figure()
w = array([2, 1])
b = 0;

drawplane(w, b, '$\mathbf{w}$', 'b', 'r')

# a few points
plt.plot([2.2,4,3], [-2,2,-1], 'bx')
plt.plot([-2.2,-4,-3], [2,1,-1], 'ro')

plt.text(-1.8,4, "$f(\mathbf{x}) = 0$", fontsize=12)
plt.text(4, 3.5, "$f(\mathbf{x})>0$\nclass +1", ha="right", fontsize=12)
plt.text(-4, -3.5, "$f(\mathbf{x})<0$\nclass -1", fontsize=12)

plt.xlabel('feature $x_1$'); plt.ylabel('feature $x_2$')
plt.axis('equal')
plt.axis([-5, 5, -5, 5]); plt.grid(True)
plt.close()

```

```


<>:5: SyntaxWarning: invalid escape sequence '\m'
<>:11: SyntaxWarning: invalid escape sequence '\m'
<>:12: SyntaxWarning: invalid escape sequence '\m'
<>:13: SyntaxWarning: invalid escape sequence '\m'
<>:5: SyntaxWarning: invalid escape sequence '\m'
<>:11: SyntaxWarning: invalid escape sequence '\m'
<>:12: SyntaxWarning: invalid escape sequence '\m'
<>:13: SyntaxWarning: invalid escape sequence '\m'
/var/folders/5h/9z8cfrvj02324js7fw73_cb00000gn/T/ipykernel_43280/1863615804.py:5
: SyntaxWarning: invalid escape sequence '\m'
    drawplane(w, b, '$\mathbf{w}$', 'b', 'r')
/var/folders/5h/9z8cfrvj02324js7fw73_cb00000gn/T/ipykernel_43280/1863615804.py:1
1: SyntaxWarning: invalid escape sequence '\m'
    plt.text(-1.8,4, "$f(\mathbf{x}) = 0$", fontsize=12)
/var/folders/5h/9z8cfrvj02324js7fw73_cb00000gn/T/ipykernel_43280/1863615804.py:1
2: SyntaxWarning: invalid escape sequence '\m'
    plt.text(4, 3.5, "$f(\mathbf{x})>0$\nclass +1", ha="right", fontsize=12)
/var/folders/5h/9z8cfrvj02324js7fw73_cb00000gn/T/ipykernel_43280/1863615804.py:1
3: SyntaxWarning: invalid escape sequence '\m'
    plt.text(-4, -3.5, "$f(\mathbf{x})<0$\nclass -1", fontsize=12)

```

```
[11]: linclass
```

Ignoring fixed x limits to fulfill fixed data aspect with adjustable data limits.

```
[11]:
```



8 Separating Hyperplane

- In a d -dimensional feature space, the parameters are $\mathbf{w} \in \mathbb{R}^d$.
- The equation $\mathbf{w}^T \mathbf{x} + b = 0$ defines a $(d - 1)$ -dim. linear surface:
 - for $d = 2$, $\mathbf{w}$ defines a 1-D line.
 - for $d = 3$, $\mathbf{w}$ defines a 2-D plane.
 - ...
 - in general, we call it a hyperplane.

9 Learning the classifier

- How to set the classifier parameters $(\mathbf{w}, b)$?
 - Learn them from training data!
- Classifiers differ in the objectives used to learn the parameters $(\mathbf{w}, b)$.
 - We will look at two examples:
 - * *logistic regression*

* support vector machine (SVM)

10 Logistic regression

- Use a probabilistic approach
 - Map the linear function $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$ to probability values between 0 and 1 using a *sigmoid* function.
 - $\sigma(z) = \frac{1}{1+e^{-z}}$

```
[12]: z = linspace(-5,5)
      sigz = 1/(1+exp(-z))

      sigmoidplot = plt.figure(figsize=(5,3))


      plt.plot(z,sigz, 'b-')

      plt.xlabel('$f(\mathbf{x})$'); plt.ylabel('$\sigma(f(\mathbf{x}))$')
      plt.axis([-5, 5, 0, 1]); plt.grid(True)
      plt.close()
```

```
<>:8: SyntaxWarning: invalid escape sequence '\m'
<>:8: SyntaxWarning: invalid escape sequence '\s'
<>:8: SyntaxWarning: invalid escape sequence '\m'
<>:8: SyntaxWarning: invalid escape sequence '\s'
/var/folders/5h/9z8cfrvj02324js7fw73_cb00000gn/T/ipykernel_43280/3647619821.py:8
: SyntaxWarning: invalid escape sequence '\m'
  plt.xlabel('$f(\mathbf{x})$'); plt.ylabel('$\sigma(f(\mathbf{x}))$')
/var/folders/5h/9z8cfrvj02324js7fw73_cb00000gn/T/ipykernel_43280/3647619821.py:8
: SyntaxWarning: invalid escape sequence '\s'
  plt.xlabel('$f(\mathbf{x})$'); plt.ylabel('$\sigma(f(\mathbf{x}))$')
```

```
[13]: sigmoidplot
```

```
[13]:
```



- Given a feature vector x , the probability of a class is:
 - $p(y = +1|\mathbf{x}) = \sigma(f(\mathbf{x}))$
 - $p(y = -1|\mathbf{x}) = 1 - \sigma(f(\mathbf{x}))$
- Note: here we are directly modeling the class posterior probability!
 - not the class-conditional $p(\mathbf{x}|y)$

```
[14]: x = linspace(-10,10, 100)
w = 2.0
b = -4.0
f = w*x+b
sf = 1/(1+exp(-f))
midx = -b/w

lrexample = plt.figure(figsize=(6,3))

plt.plot(x,sf, 'b-', label="$p(y=+1|\mathbf{x})$")
plt.plot(x,1-sf, 'r-', label="$p(y=-1|\mathbf{x})$")
plt.plot([midx, midx], [0.0,1.2], 'k--', label="decision boundary")

plt.arrow(midx-0.1,1.05,-1.8,0,width=0.01)
plt.arrow(midx+0.1,1.05,1.8,0,width=0.01)
plt.text(midx+0.2,1.10,"Class +1")
plt.text(midx-0.2,1.10,"Class -1", horizontalalignment='right')
plt.legend(loc=0, framealpha=1, fontsize='medium')
plt.title('class posterior $p(y|\mathbf{x})$');
plt.xlabel('feature $\mathbf{x}$'); plt.ylabel('probability')
plt.axis([-7, 7, 0, 1.2]); plt.grid(True)
plt.close()
```


```

<>:10: SyntaxWarning: invalid escape sequence '\m'
<>:11: SyntaxWarning: invalid escape sequence '\m'
<>:19: SyntaxWarning: invalid escape sequence '\m'
<>:20: SyntaxWarning: invalid escape sequence '\m'
<>:10: SyntaxWarning: invalid escape sequence '\m'
<>:11: SyntaxWarning: invalid escape sequence '\m'
<>:19: SyntaxWarning: invalid escape sequence '\m'
<>:20: SyntaxWarning: invalid escape sequence '\m'
/var/folders/5h/9z8cfrvj02324js7fw73_cb00000gn/T/ipykernel_43280/3055418326.py:1
0: SyntaxWarning: invalid escape sequence '\m'
    plt.plot(x,sf, 'b-', label="$p(y=+1|\mathbf{x})$")
/var/folders/5h/9z8cfrvj02324js7fw73_cb00000gn/T/ipykernel_43280/3055418326.py:1
1: SyntaxWarning: invalid escape sequence '\m'
    plt.plot(x,1-sf, 'r-', label="$p(y=-1|\mathbf{x})$")
/var/folders/5h/9z8cfrvj02324js7fw73_cb00000gn/T/ipykernel_43280/3055418326.py:1
9: SyntaxWarning: invalid escape sequence '\m'
    plt.title('class posterior $p(y|\mathbf{x})$');
/var/folders/5h/9z8cfrvj02324js7fw73_cb00000gn/T/ipykernel_43280/3055418326.py:2
0: SyntaxWarning: invalid escape sequence '\m'
    plt.xlabel('feature $\mathbf{x}$'); plt.ylabel('probability')

```

[15]: lrexample

[15]:



11 Learning the parameters

- Given training data $\{\mathbf{x}_i, y_i\}_{i=1}^N$, learn the function parameters $(\mathbf{w}, b)$ using maximum likelihood estimation.

- maximize the likelihood of the data $\{\mathbf{x}_i, y_i\}$ according to the posterior:

$$(\mathbf{w}^*, b^*) = \underset{\mathbf{w}, b}{\operatorname{argmax}} \sum_{i=1}^N \log p(y_i | \mathbf{x}_i)$$

- posterior is a Bernoulli distribution (given $\mathbf{x}$):

$$p(y | \mathbf{x}) = \begin{cases} \sigma(f(\mathbf{x})), & y = 1 \\ 1 - \sigma(f(\mathbf{x})), & y = -1 \end{cases}$$

- Note the following property:

$$1 - \sigma(z) = \sigma(-z)$$

```
[16]: z = linspace(-5,5)
      sigz = 1/(1+exp(-z))
      signz = 1/(1+exp(z))

      sigmoidplot = plt.figure(figsize=(5,3))


      plt.plot(z,sigz, 'b-', label='$\sigma(z)$')
      plt.plot(z,signz, 'r-', label='$\sigma(-z) = 1-\sigma(z)$')

      plt.xlabel('$f(\mathbf{x})$');
      plt.axis([-5, 5, 0, 1]); plt.grid(True)
      plt.legend()
      plt.close()
```

```
<>:7: SyntaxWarning: invalid escape sequence '\s'
<>:8: SyntaxWarning: invalid escape sequence '\s'
<>:10: SyntaxWarning: invalid escape sequence '\m'
<>:7: SyntaxWarning: invalid escape sequence '\s'
<>:8: SyntaxWarning: invalid escape sequence '\s'
<>:10: SyntaxWarning: invalid escape sequence '\m'
/var/folders/5h/9z8cfrvj02324js7fw73_cb00000gn/T/ipykernel_43280/3063172703.py:7
: SyntaxWarning: invalid escape sequence '\s'
  plt.plot(z,sigz, 'b-', label='$\sigma(z)$')
/var/folders/5h/9z8cfrvj02324js7fw73_cb00000gn/T/ipykernel_43280/3063172703.py:8
: SyntaxWarning: invalid escape sequence '\s'
  plt.plot(z,signz, 'r-', label='$\sigma(-z) = 1-\sigma(z)$')
/var/folders/5h/9z8cfrvj02324js7fw73_cb00000gn/T/ipykernel_43280/3063172703.py:1
0: SyntaxWarning: invalid escape sequence '\m'
  plt.xlabel('$f(\mathbf{x})$');
```

```
[17]: sigmoidplot
```

```
[17]:
```



- Thus,

$$p(y|\mathbf{x}) = \begin{cases} \sigma(f(\mathbf{x})), & y = 1 \\ \sigma(-f(\mathbf{x})), & y = -1 \end{cases}$$

- Simplifying the 2 cases into one equation,

$$p(y|\mathbf{x}) = \sigma(yf(\mathbf{x}))$$

- Taking the log,

$$\begin{aligned} \log p(y|\mathbf{x}) &= \log \sigma(yf(\mathbf{x})) \\ &= \log \frac{1}{1 + e^{-yf(\mathbf{x})}} \\ &= -\log(1 + e^{-yf(\mathbf{x})}) \end{aligned}$$

- Substituting into the MLE formulation:

$$\begin{aligned} (\mathbf{w}^*, b^*) &= \operatorname{argmax}_{\mathbf{w}, b} \sum_{i=1}^N \log p(y_i|\mathbf{x}_i) \\ &= \operatorname{argmin}_{\mathbf{w}, b} \sum_{i=1}^N \log(1 + e^{-y_i f(\mathbf{x}_i)}) \end{aligned}$$


- the term on the right is a *data-fit term*
 - wants to make the parameters $(\mathbf{w}, b)$ to well fit the data.
 - Define $z_i = y_i f(\mathbf{x}_i)$
 - * Interesting observation:
 - $z_i > 0$ when sample $\mathbf{x}_i$ is classified correctly
 - $z_i < 0$ when sample $\mathbf{x}_i$ is classified incorrectly
 - $z_i = 0$ when sample is on classifier boundary
 - logistic loss function: $L(z_i) = \log(1 + \exp(-z_i))$

```
[18]: z = linspace(-6,6,100)
logloss = log(1+exp(-z)) / log(2)
lossfig = plt.figure()
plt.plot(z,logloss, 'b-')
plt.plot([0,0], [0,9], 'k--')
plt.text(0,8.5, "incorrectly classified $\Leftarrow$ ", ha='right',
        weight='bold')
plt.text(0,8.5, " $\Rightarrow$ correctly classified", ha='left', weight='bold')
plt.annotate(text="large loss for badly\nmisclassified samples",
             xy=(-4,5.8), xytext=(-4.5,7.2),backgroundcolor='white',
             arrowprops=dict(arrowstyle="->"))
plt.annotate(text="non-zero loss for correctly\nclassified samples\nthat are",
             xy=(1,0.5), xytext=(0.5,3.2), backgroundcolor='white',
             arrowprops=dict(arrowstyle="->"))
plt.xlabel('$z_i$'); plt.ylabel('loss')
plt.title('logistic loss'); plt.grid(True)
plt.xlim(-6,6)
plt.close()
```

```
<>:7: SyntaxWarning: invalid escape sequence '\R'
<>:7: SyntaxWarning: invalid escape sequence '\R'
/var/folders/5h/9z8cfrvj02324js7fw73_cb00000gn/T/ipykernel_43280/703837359.py:7:
SyntaxWarning: invalid escape sequence '\R'
    plt.text(0,8.5, " $\Rightarrow$ correctly classified", ha='left',
weight='bold')
```

```
[19]: lossfig
```

```
[19]:
```



12 Regularization

- to prevent *overfitting*, add a prior distribution on $\mathbf{w}$.
 - prefer solutions that are likely under the prior.

$$(\mathbf{w}^*, b^*) = \underset{\mathbf{w}, b}{\operatorname{argmax}} \log p(\mathbf{w}) + \sum_{i=1}^N \log p(y_i | \mathbf{x}_i)$$

- assume Gaussian distribution on $\mathbf{w}$ with variance $C/2$
 - $p(\mathbf{w}) = N(\mathbf{w} | 0, \frac{C}{2} \mathbf{I})$
 - * small values of C keep $\mathbf{w}$ close to 0.
 - * large values of C allow larger values of $\mathbf{w}$.
 - $\log p(\mathbf{w}) = -\frac{1}{C} \mathbf{w}^T \mathbf{w} + \text{constant}$

- Substituting,

$$(\mathbf{w}^*, b^*) = \underset{\mathbf{w}, b}{\operatorname{argmin}} \frac{1}{C} \mathbf{w}^T \mathbf{w} + \sum_{i=1}^N \log(1 + \exp(-y_i f(\mathbf{x}_i)))$$

- the first term is the *regularization term*
 - Note: $\mathbf{w}^T \mathbf{w} = \sum_{j=1}^d w_j^2$
 - penalty term that keeps entries in $\mathbf{w}$ from getting too large.

- C is the regularization *hyperparameter*
 - * larger C values apply less penalty on large $\mathbf{w}$ → allow large values in $\mathbf{w}$.
 - * smaller C values apply more penalty on large $\mathbf{w}$ → discourage large values in $\mathbf{w}$.
- the second term is the *data fit term* - same as before.

13 Optimization

- **no closed-form solution**
 - use an iterative optimization algorithm to find the optimal solution
 - e.g., *gradient descent* - step downhill in each iteration.
 - * $\mathbf{w} \leftarrow \mathbf{w} - \eta \frac{dE}{d\mathbf{w}}$
 - * where E is the objective function
 - * η is the *learning rate* (how far to step in each iteration).

14 Example: Iris Data

```
[20]: # load iris data each row is (petal length, sepal width, class)
irisdata = loadtxt('iris2.csv', delimiter=',', skiprows=1)

X = irisdata[:,0:2] # the first two columns are features (petal length, sepal
↪width)
Y = irisdata[:,2]   # the third column is the class label (versicolor=1,
↪virginica=2)
                        # --> automatically mapped to (-1, +1) when training
↪classifier

print(X.shape)
```


(100, 2)

```
[21]: # a colormap for making the scatter plot: class -1 will be red, class +1 will
↪be green
mycmap = matplotlib.colors.LinearSegmentedColormap.from_list('mycmap',
↪["#FF0000", "#FFFFFF", "#00FF00"])

axbox = [2.5, 7, 1.5, 4] # common axis range

# a function for setting a common plot
def irisaxis(axbox):
    plt.xlabel('petal length'); plt.ylabel('sepal width')
    plt.axis(axbox); plt.grid(True)
```

```
[22]: # show the data
plt.figure()
plt.scatter(X[:,0], X[:,1], c=Y, cmap=mycmap, edgecolors='k')
irisaxis(axbox)
```



```
[23]: # randomly split data into 50% train and 50% test set
trainX, testX, trainY, testY = \
    model_selection.train_test_split(X, Y,
    train_size=0.5, test_size=0.5, random_state=4487)

print(trainX.shape)
print(testX.shape)
```

```
(50, 2)
```

```
(50, 2)
```

```
[24]: # learn logistic regression classifier
# (C is a regularization hyperparameter)
logreg = linear_model.LogisticRegression(C=100)
logreg.fit(trainX, trainY)

print("w =", logreg.coef_)
print("b =", logreg.intercept_)
```

```
w = [[9.53947455 0.8988902 ]]
```

```
b = [-48.82195461]
```

- Equation:
 - $f(\mathbf{x}) = (9.51 * \text{petal_length}) + (0.895 * \text{sepal_width}) - 48.68$
- Interpretation:
 - large petal length makes $f(\mathbf{x})$ positive, so large petal length is associated with class +1.


```
[25]: def plot_posterior(model, axbox, mycmap, showlabels=True):
    xr = [ linspace(axbox[0], axbox[1], 200),
           linspace(axbox[2], axbox[3], 200) ]

    # make a grid for calculating the posterior,
    # then form into a big [N,2] matrix
    xgrid0, xgrid1 = meshgrid(xr[0], xr[1])
    allpts = c_[xgrid0.ravel(), xgrid1.ravel()]

    # calculate the posterior probability
    post = model.predict_proba(allpts)
    # extract the posterior for class 2, and reshape into a grid
    post1 = post[:,1].reshape(xgrid0.shape)

    # contour plot of the posterior and decision boundary
    plt.imshow(post1, origin='lower', extent=axbox, alpha=0.50, cmap=mycmap,
    ↪vmin=0.0, vmax=1.0)
    if showlabels:
        plt.colorbar(shrink=0.6)
        CS = plt.contour(xr[0], xr[1], post1, colors='k', linestyle='dashed',
    ↪levels=[0.1, 0.3, 0.7, 0.9], alpha=0.8, linewidths=1)
        if showlabels:
            #plt.colorbar(CS)
            plt.clabel(CS, inline=1, fontsize=10)
        plt.contour(xr[0], xr[1], post1, levels=[0.5], linewidths=2, colors='black')
        irisaxis(axbox)
```

```
[26]: # show the posterior and training data
plt.figure(figsize=(8,6))
plot_posterior(logreg, axbox, mycmap)
plt.scatter(trainX[:,0], trainX[:,1], c=trainY, cmap=mycmap, edgecolors='k')
plt.title('posterior p(y=+1|x) with training data');
```



```
[27]: # predict from the model
predY = logreg.predict(testX)


# calculate accuracy
acc = metrics.accuracy_score(testY, predY)
print("test accuracy =", acc)
```

test accuracy = 0.92

```
[28]: # show the posterior and training data
postfig = plt.figure(figsize=(8,6))
plot_posterior(logreg, axbox, mycmap)
plt.scatter(trainX[:,0], trainX[:,1], c=trainY, cmap=mycmap, marker="o",
            label="training", edgecolors='k')
plt.scatter(testX[:,0], testX[:,1], c=testY, cmap=mycmap, marker="s",
            label="testing", edgecolors='k')
plt.title('posterior p(y=+1|x) with testing data');
plt.legend(loc=0);
plt.close()
```

```
[29]: postfig
```

```
[29]:
```



15 Selecting the regularization hyperparameter

- the regularization hyperparameter C has a large effect on the decision boundary and the accuracy.
 - larger C makes the classifier more confident (posterior probabilities saturate to 0 and 1)
 - * more likely to overfit
 - smaller C makes the classifier less confident (wider range of posterior probabilities).
 - * less likely to overfit
- How to set the value of C ?

```
[30]: lrC = plt.figure(figsize=(10,4.5))

allC = [10000,100, 10, 1, 0.1, 0.01]
for (myCind,myC) in enumerate(allC):
    logreg = linear_model.LogisticRegression(C=myC)
    logreg.fit(trainX, trainY)

    # predict from the model
    predY = logreg.predict(testX)


    # calculate accuracy
    acc = metrics.accuracy_score(testY, predY)

    plt.subplot(2,3,myCind+1)
    plot_posterior(logreg, axbox, mycmap, showlabels=False)
    plt.scatter(trainX[:,0], trainX[:,1], c=trainY, cmap=mycmap, edgecolors='k')
```

```
plt.title('C='+str(myC)+"; acc="+str(acc));
plt.gca().xaxis.set_ticklabels([])
plt.gca().yaxis.set_ticklabels([])
plt.close()
```

[31]: lrC

[31]:



16 Cross-validation

- Use *cross-validation* on the training set to select the best value of C .
 - Run many experiments on the training set to see which parameters work on different versions of the data.
 - * Split the data into batches of training and validation data.
 - * Try a range of C values on each split.
 - * Pick the value that works best over all splits.
- **Procedure**
 1. select a range of C values to try
 2. Repeat K times
 3. Split the training set into training data and validation data
 4. Learn a classifier for each value of C
 5. Record the accuracy on the validation data for each C
 6. Select the value of C that has the highest average accuracy over all K folds.
 7. Retrain the classifier using all data and the selected C .
- scikit-learn already has built-in `cross_validation` module (more later).
- for logistic regression, use `LogisticRegressionCV` class

```
[32]: # learn logistic regression classifier using CV
# Cs is an array of possible C values
```

```

# cv is the number of folds
# n_jobs=-1 means run in parallel with all cores
logreg = linear_model.LogisticRegressionCV(Cs=logspace(-4,4,20), cv=5,
    ↪n_jobs=-1)
logreg.fit(trainX, trainY)

print("w=", logreg.coef_)
print("b=", logreg.intercept_)

# predict from the model
predY = logreg.predict(testX)

# calculate accuracy
acc = metrics.accuracy_score(testY, predY)
print("test accuracy=", acc)

```

```

w= [[4.62056477 0.72461222]]
b= [-24.25601763]
test accuracy= 0.9

```

17 Which C was selected?

```


[33]: print("C =", logreg.C_)
# calculate the average score for each C
avgscores = mean(logreg.scores_[2],0) # 2 is the class label
plt.semilogx(logreg.Cs_, avgscores, 'ko-')
plt.xlabel('C'); plt.ylabel('average CV accuracy'); plt.grid(True);

```

```

C = [4.2813324]

```



18 Multi-class classification

- So far, we have only learned a classifier for 2 classes (+1, -1)
 - called a **binary classifier**
- For more than 2 classes, split the problem up into several binary classifier problems.
 - **1-vs-rest**
 - * *Training*: for each class, train a classifier for that class versus the other classes.
 - For example, if there are 3 classes, then train 3 binary classifiers: 1 vs {2,3}; 2 vs {1,3}; 3 vs {1,2}
 - * *Prediction*: calculate probability for each binary classifier. Select the class with highest probability.

19 Example on 3-class Iris data

```
[34]: # load iris data each row is (petal length, sepal width, class)
irisdata = loadtxt('iris3.csv', delimiter=',', skiprows=1)

X = irisdata[:,0:2] # the first two columns are features (petal length, sepal
↪width)
```

```
Y = irisdata[:,2]    # the third column is the class label (setosa=0,
↪versicolor=1, virginica=2)

print(X.shape)
```

(150, 2)

```
[35]: # randomly split data into 50% train and 50% test set
trainX, testX, trainY, testY = \
    model_selection.train_test_split(X, Y,
    train_size=0.5, test_size=0.5, random_state=4487)

print(trainX.shape)
print(testX.shape)
```


(75, 2)

(75, 2)

```
[36]: # look at training data

axbox3 = [0.8, 7, 1.5, 4.5]
# make a colormap for viewing 3 classes
mycmap3 = matplotlib.colors.LinearSegmentedColormap.from_list('mycmap',
↪["#FF0000", "#00FF00", "#0000FF"])

plt.scatter(trainX[:,0], trainX[:,1], c=trainY, cmap=mycmap3, edgecolors='k')
plt.axis(axbox3); plt.grid(True);
plt.xlabel('petal length'); plt.ylabel('sepal width');
```



```
[37]: # learn logistic regression classifier (one-vs-all)
mlogreg = linear_model.LogisticRegression(C=10, multi_class='ovr')
mlogreg.fit(trainX, trainY)

# now contains 3 hyperplanes and 3 bias terms (one for each class)
print("w=", mlogreg.coef_)
print("b=", mlogreg.intercept_)

# predict from the model
predY = mlogreg.predict(testX)

# calculate accuracy
acc = metrics.accuracy_score(testY, predY)
print("test accuracy=", acc)
```

```
w= [[-3.56993402  1.08723834]
     [-0.03184394 -2.23108274]
     [ 5.42003067 -1.68985998]]
b= [ 6.41058825  6.12262666 -21.81518434]
test accuracy= 0.9733333333333334
```

```
/Users/zzs/miniconda3/envs/cs5489/lib/python3.12/site-
```

packages/sklearn/linear_model/_logistic.py:1281: FutureWarning: 'multi_class' was deprecated in version 1.5 and will be removed in 1.7. Use OneVsRestClassifier(LogisticRegression(...)) instead. Leave it to its default value to avoid this warning.

warnings.warn(

```
[38]: def plot_1vr_classifiers(logreg, mlogreg, axbox, mycmap, trainX, trainY,
    ↪titstr):
    # plot each classifier (assume 3)
    for i in range(3):
        plt.subplot(1,3,i+1)
        # make binary model
        ilogreg = linear_model.LogisticRegression(logreg)
        ilogreg.coef_ = mlogreg.coef_[i,:].reshape(1,2)
        ilogreg.intercept_ = mlogreg.intercept_[i].reshape(1,)
        ilogreg.classes_ = array([1, 2])
        itrainY = (trainY == i)

        plot_posterior(ilogreg, axbox, mycmap, showlabels=False)
        plt.scatter(trainX[:,0], trainX[:,1], c=itrainY, cmap=mycmap,
    ↪edgecolors='k')

        plt.gca().axis.set_ticklabels([])
        plt.gca().yaxis.set_ticklabels([])
        plt.title(titstr.format(i))

    mlrfig = plt.figure(figsize=(9,6))
    plot_1vr_classifiers(logreg, mlogreg, axbox3, mycmap, trainX, trainY, "class {}
    ↪vs. rest")
    plt.close()
```


- the individual 1-vs-rest binary classifiers

```
[39]: print("w=", mlogreg.coef_)
    print("b=", mlogreg.intercept_)

    mlrfig

w= [[-3.56993402  1.08723834]
    [-0.03184394 -2.23108274]
    [ 5.42003067 -1.68985998]]
b= [ 6.41058825  6.12262666 -21.81518434]
```

[39]:



```
[40]: def plot_posterior3(model, axbox, mycmap):
    xr = [ arange(0.8,7,0.05) , arange(1.5, 4.5, 0.05) ]

    # make a grid for calculating the posterior,
    # then form into a big [N,2] matrix
    xgrid0, xgrid1 = meshgrid(xr[0], xr[1])
    allpts = c_[xgrid0.ravel(), xgrid1.ravel()]

    # predict probabilities
    Z = model.predict_proba(allpts)
    P = model.predict(allpts)

    # use probabilities as RGB color
    ZZ = Z.reshape((len(xr[1]), len(xr[0]), 3))


    plt.imshow(ZZ, origin='lower', extent=axbox, alpha=0.50)
    plt.contour(xr[0], xr[1], P.reshape(xgrid0.shape), levels=[0.5,1.5,2.5],
    ↳linestyles='dashed', colors='black')
    irisaxis(axbox)

    lr3class = plt.figure()
    plot_posterior3(mlogreg, axbox3, mycmap3)
    plt.scatter(trainX[:,0], trainX[:,1], c=trainY, cmap=mycmap3, edgecolors='k')
    plt.title('class probabilities');
    plt.close()
```

- the final classifier, combining all 1 vs rest classifiers

```
[41]: lr3class
```

```
[41]:
```



20 Multiclass logistic regression

- Another way to get a multi-class classifier is to define a multi-class objective.
 - One weight vector $\mathbf{w}_c$ for each class c .
 - linear function for each class, $f_c(\mathbf{x}) = \mathbf{w}_c^T \mathbf{x}$.
- Define probabilities with **softmax** function
 - analogous to sigmoid function for binary logistic regression.

$$p(y = c|\mathbf{x}) = \frac{e^{f_c(\mathbf{x})}}{e^{f_1(\mathbf{x})} + \dots + e^{f_K(\mathbf{x})}}$$

- The class with largest response of $f_c(\mathbf{x})$ will have the highest probability.

```
[42]: f1 = linspace(-5,5,101)
f2 = linspace(-5,5,101)

sfmfig = plt.figure(figsize=(10,4))
for i,f3 in enumerate([-1,3]):
    xgrid0, xgrid1 = meshgrid(f1, f2)
    plt.subplot(1,2,i+1)
    p = exp(xgrid0) / (exp(xgrid0) + exp(xgrid1)+exp(f3))

    plt.imshow(p, origin='lower', extent=[-5,5,-5,5], vmin=0, vmax=1)
    plt.colorbar(shrink=0.8)
    CS = plt.contour(f1, f2, p, colors='k', linestyle='dashed', levels=[0.1, 0.
↪3, 0.5, 0.7, 0.9], alpha=0.8, linewidths=1)

    plt.clabel(CS, inline=1, fontsize=10)
```

```
plt.xlabel('$f_1$')
plt.ylabel('$f_2$')
plt.title('$p(y=1|\mathbf{x})$, $f_3=$' + str(f3) + '$')
plt.close()
```


```
<>:18: SyntaxWarning: invalid escape sequence '\m'
<>:18: SyntaxWarning: invalid escape sequence '\m'
/var/folders/5h/9z8cfrvj02324js7fw73_cb00000gn/T/ipykernel_43280/696709212.py:18
: SyntaxWarning: invalid escape sequence '\m'
plt.title('$p(y=1|\mathbf{x})$, $f_3=$' + str(f3) + '$')
```

- Example with $K = 3$.

$$p(y = 1|\mathbf{x}) = \frac{e^{f_1(\mathbf{x})}}{e^{f_1(\mathbf{x})} + e^{f_2(\mathbf{x})} + e^{f_3(\mathbf{x})}}$$

[43]: sfmfig

[43]:



21 Parameter estimation

- Estimate the $\{\mathbf{w}_j\}$ parameters using MLE.
- Let $(\mathbf{x}, \mathbf{y})$ be a data sample pair:
 - $\mathbf{x}$ feature vector.
 - $\mathbf{y} = [y_1, \dots, y_K]$ is a one-hot vector, where $y_c = 1$ when class c , and 0 otherwise.

- Data likelihood of $(\mathbf{x}, \mathbf{y})$.

$$\text{likelihood:} \quad p(\mathbf{y}|\mathbf{x}) = \prod_{j=1}^K p(y = j|\mathbf{x})^{y_j}$$

$$\text{log-likelihood:} \quad \log p(\mathbf{y}|\mathbf{x}) = \sum_{j=1}^K y_j \log p(y = j|\mathbf{x})$$

$$\text{negative log-likelihood:} \quad -\log p(\mathbf{y}|\mathbf{x}) = -\sum_{j=1}^K y_j \log p(y = j|\mathbf{x})$$

– equivalent to the *cross-entropy loss*

- Given dataset $\{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^N$
 - maximize the data log-likelihood:

$$\max_{\{\mathbf{w}_j\}} \sum_{i=1}^N \log p(\mathbf{y}_i|\mathbf{x}_i) = \max_{\{\mathbf{w}_j\}} \sum_{i=1}^N \sum_{j=1}^K y_{ij} \log p(y = j|\mathbf{x}_i)$$

– i.e., minimize the cross-entropy loss

```
[44]: # learn logistic regression classifier
mlogreg = linear_model.LogisticRegression(C=10,
                                          multi_class='multinomial')
      # use multi-class and corresponding solver
mlogreg.fit(trainX, trainY)

# now contains 3 hyperplanes and 3 bias terms (one for each class)
print("w=", mlogreg.coef_)
print("b=", mlogreg.intercept_)

# predict from the model
predY = mlogreg.predict(testX)

# calculate accuracy
acc = metrics.accuracy_score(testY, predY)
print("test accuracy=", acc)
```


```
w= [[-4.13160888  1.30538269]
     [-0.71650483  0.23668666]
     [ 4.84811371 -1.54206935]]
b= [ 11.46619303   5.40363007 -16.8698231 ]
test accuracy= 0.9733333333333334
```

```
/Users/zzs/miniconda3/envs/cs5489/lib/python3.12/site-
packages/sklearn/linear_model/_logistic.py:1272: FutureWarning: 'multi_class'
was deprecated in version 1.5 and will be removed in 1.7. From then on, it will
always use 'multinomial'. Leave it to its default value to avoid this warning.
warnings.warn(
```

```
[45]: lr3classm = plt.figure()
plot_posterior3(mlogreg, axbox3, mycmap3)
plt.scatter(trainX[:,0], trainX[:,1], c=trainY, cmap=mycmap3, edgecolors='k')
plt.title('class probabilities');
plt.close()
```

```
[46]: lr3classm
```

[46]:



```
[47]: lr31vr = plt.figure(figsize=(9,6))
plot_1vr_classifiers(logreg, mlogreg, axbox3, mycmap, trainX, trainY, "w-{}")
plt.close()
```


- individual weight vectors work together to partition the space

```
[48]: print("w=", mlogreg.coef_)
print("b=", mlogreg.intercept_)
```

```
lr31vr
```

```
w= [[-4.13160888  1.30538269]
     [-0.71650483  0.23668666]
     [ 4.84811371 -1.54206935]]
b= [ 11.46619303   5.40363007 -16.8698231 ]
```

[48]:



22 Logistic Regression Summary

- **Classifier:**
 - linear function $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$
 - Given a feature vector $\mathbf{x}$, the probability of a class is:
 - * $p(y = +1|\mathbf{x}) = \sigma(f(\mathbf{x}))$
 - * $p(y = -1|\mathbf{x}) = 1 - \sigma(f(\mathbf{x}))$
 - * *sigmoid* function: $\sigma(z) = \frac{1}{1+e^{-z}}$
 - logistic loss function: $L(z) = \log(1 + \exp(-z))$
- **Training:**
 - Maximize the likelihood of the training data.
 - Use regularization to prevent overfitting.
 - * Use cross-validation to pick the regularization hyperparameter C .
- **Classification:**
 - Given a new sample $\mathbf{x}^*$:
 - * pick class with highest probability $p(y|\mathbf{x}^*)$:

$$y^* = \begin{cases} +1, & p(y = +1|\mathbf{x}^*) > p(y = -1|\mathbf{x}^*) \\ -1, & \text{otherwise} \end{cases}$$

- * alternatively, just use $f(\mathbf{x}^*)$

$$y^* = \begin{cases} +1, & f(\mathbf{x}^*) > 0 \\ -1, & \text{otherwise} \end{cases} = \text{sign}(f(\mathbf{x}_*))$$

- **Extend to multi-class:**
 - K linear functions, one for each class.
 - compute probability using softmax function
 - MLE equivalent to cross-entropy loss