

Mechanistic Interpretability of LLM Reasoning via Knowledge Manipulation

Interim Report

HAN Zhuo (Student ID: 60130387)

March 10, 2026

1 Introduction

Large language models (LLMs) have demonstrated strong performance on a wide range of reasoning tasks [32, 3, 24], including multi-step question answering, symbolic inference, and rule-based decision making [4, 22]. Their success has raised an important question: how do these models internally perform reasoning? Although recent work has shown that LLMs can generate impressive reasoning outputs, most evidence remains behavioral, typically based on final-answer accuracy or the quality of generated reasoning traces [13, 34, 28]. Such results indicate that reasoning-like behavior is present, but they do not by themselves explain how reasoning is implemented inside the model.

This limitation has made interpretability an increasingly important direction for reasoning research. In particular, mechanistic interpretability seeks to explain model behavior in terms of internal computational components and their functional interactions [8, 6, 14, 30, 16]. Rather than treating LLMs as black boxes, it aims to identify how layers, attention heads, neurons, or larger circuits contribute to specific behaviors. For reasoning tasks, this perspective is especially valuable because successful performance depends not only on producing the correct answer, but also on how relevant information is retrieved, maintained, transformed, and combined across intermediate computational steps.

However, current approaches to reasoning interpretability still face a central challenge. Many existing studies rely on observational evidence, such as attention patterns, attribution scores, or probing-based signals [29, 19, 33]. Although these methods can reveal useful correlations, they often provide limited evidence about functional necessity. An internal feature may co-occur with successful reasoning without being causally responsible for it. As a result, a more intervention-based framework is needed in order to move from descriptive analysis toward mechanistic explanation.

This paper proposes to address this challenge through knowledge manipulation. The core idea is to treat knowledge not merely as a static capability stored in the model, but as an experimental variable for probing reasoning mechanisms. By introducing controlled changes to the knowledge conditions underlying a reasoning task, it becomes possible to observe how these changes influence both model outputs and internal computations. If a particular internal component genuinely supports knowledge-dependent reasoning, then targeted modifications to relevant knowledge should induce systematic changes in that component’s behavior. In this way, knowledge manipulation functions as a causal probe for identifying how reasoning is realized inside the model.

Building on this perspective, the present study investigates how controlled changes in knowledge conditions affect internal reasoning processes in LLMs. Rather than focusing only on final-

answer correctness, the study aims to trace how manipulated knowledge propagates through internal representations and contributes to reasoning outcomes. This provides a more mechanistically grounded perspective on LLM reasoning and connects reasoning research with intervention-based interpretability.

2 Background

2.1 Reasoning in Large Language Models

Reasoning has become one of the most actively studied capabilities of LLMs. A large body of work has shown that these models can solve tasks involving multi-step inference, compositional question answering, mathematical reasoning, and rule application [12, 35, 22]. In many cases, their outputs suggest an ability to process intermediate information and produce structured conclusions rather than simple pattern completion.

At the same time, the interpretation of such behavior remains contested. Strong task performance does not necessarily imply that the model is reasoning in a robust or human-like sense. More importantly, even when a model succeeds on a reasoning benchmark, the internal basis of that success often remains unclear [2, 20]. A correct answer may result from shallow shortcuts, memorized correlations, partial pattern matching, or genuinely knowledge-dependent inference [26, 36]. Behavioral evaluation alone is often insufficient to distinguish among these possibilities [11, 37].

For this reason, understanding reasoning in LLMs requires not only measuring outputs but also analyzing the internal computational processes that generate them. This shift from behavioral performance to internal mechanism is central to the present study.

2.2 Mechanistic Interpretability

Mechanistic interpretability provides a framework for analyzing model behavior at the level of internal computation [21, 1]. Its goal is to explain how specific components of a neural model contribute to the transformations that ultimately produce an output. Existing work in this area has developed methods such as activation analysis, causal tracing, activation patching, and circuit discovery, all of which aim to uncover the computational structure underlying model behavior [5, 10, 14].

This framework is particularly relevant for reasoning because reasoning is inherently process-oriented. A model solving a reasoning task must not only access relevant information, but also preserve and transform that information across multiple stages of computation [17, 27, 23]. Mechanistic interpretability makes it possible to investigate where such operations occur, how they are distributed across layers or modules, and which components are functionally important.

Nevertheless, one methodological difficulty remains prominent. Much interpretability evidence is correlational rather than causal. A component may appear important because it is active during successful reasoning, yet its activity alone does not prove that it plays a necessary role [10, 5]. To make stronger mechanistic claims, interpretability analysis benefits from controlled interventions that alter the conditions of reasoning and reveal how internal computations respond.

2.3 Knowledge Manipulation as an Interpretability Tool

A promising form of controlled intervention is knowledge manipulation. Many reasoning tasks depend on the use of relevant knowledge, whether in the form of facts, premises, rules, or structured

conditions [9, 7, 25]. If this knowledge is changed in a targeted way, the resulting differences can reveal how reasoning depends on it internally.

Knowledge manipulation is often studied in the context of knowledge editing or factual robustness [31, 14, 15, 18]. In those settings, the main question is whether the model can be updated or whether its answers remain stable under changed facts. The present study takes a different perspective. Here, knowledge manipulation is not the final objective, but a tool for understanding reasoning mechanisms. By systematically altering knowledge conditions, one can examine which internal representations are sensitive to those changes, where the effect emerges in the computation, and whether the affected components play a functional role in producing the final answer.

This perspective is useful because it links two dimensions that are often studied separately: knowledge and reasoning. In many tasks, reasoning is not independent of knowledge; it is a computation over knowledge-bearing representations. Manipulating knowledge therefore provides a principled way to probe reasoning from the inside. If a model’s internal reasoning process genuinely depends on a given premise or fact, then modifying that knowledge should lead to structured and traceable changes in internal states. Such changes can provide stronger evidence for mechanism than observation alone.

3 Methods

3.1 Study Objective

The objective of this study is to investigate the internal mechanisms of knowledge-dependent reasoning in LLMs through controlled knowledge manipulation. The analysis focuses on how targeted changes to relevant knowledge conditions affect model behavior at both the output level and the internal representation level. More specifically, the study seeks to identify which internal components are involved in retrieving, maintaining, or transforming knowledge during reasoning.

3.2 Task Design

The study focuses on reasoning tasks in which knowledge dependence is explicit and intervention points can be clearly defined. Appropriate task settings include structured multi-step reasoning, multi-hop question answering, and rule-based inference, where a correct answer depends on identifiable facts, premises, or conditions. These tasks are particularly suitable because they allow the construction of controlled example variants that differ only in a targeted knowledge-related dimension.

For each original example, manipulated counterparts are created by modifying one meaningful aspect of the knowledge condition while preserving the overall task structure. Such manipulations may include factual substitution, premise perturbation, rule-condition modification, conflict injection, or the insertion of distracting information. The purpose of this design is to isolate the effect of knowledge changes while minimizing unrelated variation across examples.

This paired construction allows the study to compare not only whether the answer changes, but also how the internal reasoning process changes when the knowledge basis of the problem is altered.

3.3 Knowledge Manipulation Design

The effectiveness of this framework depends on the design of the manipulations themselves. Each manipulation must be targeted, interpretable, and semantically controlled. The goal is not to create

arbitrary adversarial noise, but to alter a specific knowledge-bearing element that is relevant to the reasoning process.

For this reason, manipulations are designed to preserve as much of the original problem structure as possible while changing one meaningful premise, fact, or rule. This makes it easier to attribute observed differences in internal states to the manipulated knowledge condition rather than to incidental variation in wording or task form. It also makes the resulting comparisons more suitable for mechanistic interpretation.

Under this design, manipulated examples function as probes for the model’s internal reasoning process. They create controlled opportunities to test whether particular internal signals track, encode, or use the modified knowledge during inference.

3.4 Mechanistic Analysis

The core analysis examines how knowledge manipulation affects internal model states. Several interpretability methods are used for this purpose, including hidden-state comparison across original and manipulated conditions, layer-wise sensitivity analysis, and causal intervention methods such as activation patching. These tools make it possible to identify which parts of the model are most responsive to changes in knowledge and whether those parts contribute functionally to the final prediction.

The analysis is designed to move beyond output-level perturbation testing. Instead of asking only whether the answer changes, it asks where the internal computation changes, how that change propagates through the network, and whether altering a given internal signal changes the resulting reasoning behavior. This provides a stronger basis for mechanistic claims than descriptive comparison alone.

In particular, if restoring an internal activation from the original condition to the manipulated condition also restores the original reasoning behavior, that would provide evidence that the corresponding component plays a causal role in mediating the use of knowledge during reasoning. Conversely, if a component shows large observational differences but patching it has little effect, its functional importance may be limited.

3.5 Evaluation

Evaluation is conducted at multiple levels. Final-answer behavior provides one indicator of the model’s sensitivity to manipulated knowledge, but it is not sufficient on its own. The study also evaluates differences in internal representations between original and manipulated conditions, the degree to which the effects of manipulation can be localized to particular layers or components, and whether causal interventions on those components alter reasoning behavior in predictable ways.

This multi-level evaluation framework is intended to capture not only whether reasoning changes, but also how it changes internally. In doing so, it supports a more precise account of the relationship between knowledge and reasoning in LLMs.

References

- [1] BERESKA, L., AND GAVVES, E. Mechanistic interpretability for ai safety—a review. *arXiv preprint arXiv:2404.14082* (2024).
- [2] BERGLUND, L., TONG, M., KAUFMANN, M., BALESI, M., STICKLAND, A. C., KORBAK, T., AND EVANS, O. The reversal curse: LLMs trained on “a is b” fail to learn “b is a”. *arXiv preprint arXiv:2309.12288* (2023).

- [3] CHOWDHERY, A., NARANG, S., DEVLIN, J., BOSMA, M., MISHRA, G., ROBERTS, A., BARHAM, P., CHUNG, H. W., SUTTON, C., GEHRMANN, S., ET AL. Palm: Scaling language modeling with pathways. *Journal of machine learning research* 24, 240 (2023), 1–113.
- [4] CLARK, P., COWHEY, I., ETZIONI, O., KHOT, T., SABHARWAL, A., SCHOENICK, C., AND TAFJORD, O. Think you have solved question answering? try arc, the ai2 reasoning challenge. *arXiv preprint arXiv:1803.05457* (2018).
- [5] CONMY, A., MAVOR-PARKER, A., LYNCH, A., HEIMERSHEIM, S., AND GARRIGA-ALONSO, A. Towards automated circuit discovery for mechanistic interpretability. *Advances in Neural Information Processing Systems* 36 (2023), 16318–16352.
- [6] DAI, D., DONG, L., HAO, Y., SUI, Z., CHANG, B., AND WEI, F. Knowledge neurons in pretrained transformers. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (2022), pp. 8493–8502.
- [7] GEVA, M., KHASHABI, D., SEGAL, E., KHOT, T., ROTH, D., AND BERANT, J. Did aristotle use a laptop? a question answering benchmark with implicit reasoning strategies. *Transactions of the Association for Computational Linguistics* 9 (2021), 346–361.
- [8] GEVA, M., SCHUSTER, R., BERANT, J., AND LEVY, O. Transformer feed-forward layers are key-value memories. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing* (2021), pp. 5484–5495.
- [9] HAN, S., SCHOELKOPF, H., ZHAO, Y., QI, Z., RIDDELL, M., ZHOU, W., COADY, J., PENG, D., QIAO, Y., BENSON, L., ET AL. Folio: Natural language reasoning with first-order logic. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing* (2024), pp. 22017–22031.
- [10] HEIMERSHEIM, S., AND NANDA, N. How to use and interpret activation patching. *arXiv preprint arXiv:2404.15255* (2024).
- [11] INGER, N. C., ELISHA, Y., SHAPIRA, B., ROKACH, L., AND COHEN, S. Forget what you know about llms evaluations-llms are like a chameleon. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing* (2025), pp. 21675–21688.
- [12] KOJIMA, T., GU, S. S., REID, M., MATSUO, Y., AND IWASAWA, Y. Large language models are zero-shot reasoners. *Advances in neural information processing systems* 35 (2022), 22199–22213.
- [13] LYU, Q., HAVALDAR, S., STEIN, A., ZHANG, L., RAO, D., WONG, E., APIDIANAKI, M., AND CALLISON-BURCH, C. Faithful chain-of-thought reasoning. In *Proceedings of the 13th International Joint Conference on Natural Language Processing and the 3rd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)* (2023), pp. 305–329.
- [14] MENG, K., BAU, D., ANDONIAN, A., AND BELINKOV, Y. Locating and editing factual associations in gpt. *Advances in neural information processing systems* 35 (2022), 17359–17372.
- [15] MENG, K., SHARMA, A. S., ANDONIAN, A., BELINKOV, Y., AND BAU, D. Mass-editing memory in a transformer. *arXiv preprint arXiv:2210.07229* (2022).

- [16] NANDA, N., CHAN, L., LIEBERUM, T., SMITH, J., AND STEINHARDT, J. Progress measures for grokking via mechanistic interpretability. *arXiv preprint arXiv:2301.05217* (2023).
- [17] OLSSON, C., ELHAGE, N., NANDA, N., JOSEPH, N., DASARMA, N., HENIGHAN, T., MANN, B., ASKELL, A., BAI, Y., CHEN, A., ET AL. In-context learning and induction heads. *arXiv preprint arXiv:2209.11895* (2022).
- [18] PARK, H., CHOI, G., KIM, M., AND JO, Y. Context-robust knowledge editing for language models. In *Findings of the Association for Computational Linguistics: ACL 2025* (2025), pp. 10360–10385.
- [19] PAUL, D., WEST, R., BOSSELUT, A., AND FALTINGS, B. Making reasoning matter: Measuring and improving faithfulness of chain-of-thought reasoning. In *Findings of the Association for Computational Linguistics: EMNLP 2024* (2024), pp. 15012–15032.
- [20] PRESS, O., ZHANG, M., MIN, S., SCHMIDT, L., SMITH, N. A., AND LEWIS, M. Measuring and narrowing the compositionality gap in language models. In *Findings of the Association for Computational Linguistics: EMNLP 2023* (2023), pp. 5687–5711.
- [21] RAI, D., ZHOU, Y., FENG, S., SAPAROV, A., AND YAO, Z. A practical review of mechanistic interpretability for transformer-based language models. *arXiv preprint arXiv:2407.02646* (2024).
- [22] RAVICHANDER, A., NAIK, A., ROSE, C., AND HOVY, E. Equate: A benchmark evaluation framework for quantitative reasoning in natural language inference. In *Proceedings of the 23rd conference on computational natural language learning (CoNLL)* (2019), pp. 349–361.
- [23] REN, J., GUO, Q., YAN, H., LIU, D., ZHANG, Q., QIU, X., AND LIN, D. Identifying semantic induction heads to understand in-context learning. In *Findings of the Association for Computational Linguistics: ACL 2024* (2024), pp. 6916–6932.
- [24] SUZGUN, M., SCALES, N., SCHÄRLI, N., GEHRMANN, S., TAY, Y., CHUNG, H. W., CHOWDHURY, A., LE, Q., CHI, E., ZHOU, D., ET AL. Challenging big-bench tasks and whether chain-of-thought can solve them. In *Findings of the Association for Computational Linguistics: ACL 2023* (2023), pp. 13003–13051.
- [25] TAFJORD, O., DALVI, B., AND CLARK, P. Proofwriter: Generating implications, proofs, and abductive statements over natural language. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021* (2021), pp. 3621–3634.
- [26] TANG, R., KONG, D., HUANG, L., ET AL. Large language models can be lazy learners: Analyze shortcuts in in-context learning. In *Findings of the association for computational linguistics: ACL 2023* (2023), pp. 4645–4657.
- [27] TODD, E., LI, M. L., SHARMA, A. S., MUELLER, A., WALLACE, B. C., AND BAU, D. Function vectors in large language models. *arXiv preprint arXiv:2310.15213* (2023).
- [28] TURPIN, M., MICHAEL, J., PEREZ, E., AND BOWMAN, S. Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting. *Advances in Neural Information Processing Systems 36* (2023), 74952–74965.

- [29] TUTEK, M., CHALESHTORI, F. H., MARASOVIĆ, A., AND BELINKOV, Y. Measuring chain of thought faithfulness by unlearning reasoning steps. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing* (2025), pp. 9946–9971.
- [30] WANG, K., VARIENGIEN, A., CONMY, A., SHLEGERIS, B., AND STEINHARDT, J. Interpretability in the wild: a circuit for indirect object identification in gpt-2 small. *arXiv preprint arXiv:2211.00593* (2022).
- [31] WANG, S., ZHU, Y., LIU, H., ZHENG, Z., CHEN, C., AND LI, J. Knowledge editing for large language models: A survey. *ACM Computing Surveys* 57, 3 (2024), 1–37.
- [32] WEI, J., WANG, X., SCHUURMANS, D., BOSMA, M., XIA, F., CHI, E., LE, Q. V., ZHOU, D., ET AL. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems* 35 (2022), 24824–24837.
- [33] WU, C., WANG, X., LIAN, D., XIE, X., AND CHEN, E. A causality inspired framework for model interpretation. In *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining* (2023), pp. 2731–2741.
- [34] YANG, W., SUN, F., TAN, J., MA, X., CAO, Q., YIN, D., SHEN, H., AND CHENG, X. The mirage of model editing: Revisiting evaluation in the wild. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (2025), pp. 15336–15354.
- [35] YANG, Z., QI, P., ZHANG, S., BENGIO, Y., COHEN, W., SALAKHUTDINOV, R., AND MANNING, C. D. Hotpotqa: A dataset for diverse, explainable multi-hop question answering. In *Proceedings of the 2018 conference on empirical methods in natural language processing* (2018), pp. 2369–2380.
- [36] YUAN, Y., ZHAO, L., ZHANG, K., ZHENG, G., AND LIU, Q. Do llms overcome shortcut learning? an evaluation of shortcut challenges in large language models. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing* (2024), pp. 12188–12200.
- [37] ZHU, K., TU, S., JIN, Z., HOU, L., LI, J., AND ZHAO, J. Establishing trustworthy llm evaluation via shortcut neuron analysis. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (2025), pp. 3809–3822.